

SIMATS ENGINEERING

ASSESSMENT TOOL 1: Expert System Development Project

COURSE NAME: ARTIFICIAL INTELLIGENCE COURSE CODE: MLA01
AND EXPERT SYSTEMS

COURSE OUTCOME COVERED

CO 5: Construct rule-based expert systems using production rules, forward and backward chaining, and by distinguishing procedural and non-procedural paradigms across development stages. (BTL 3)

Assessment Tool Weightage: 40%

Total Marks: 100

Time: 90 Mins

No. of Questions: 1


Annexure A Expert System Development Project

Code of Conduct:

I ARTHIA (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.

I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student:



Criterion	Weightage	Marks Obtained
1. Problem Analysis and Domain Understanding	10%	
2. Knowledge Acquisition and Knowledge Representation	10%	
3. Production Rule / Knowledge Base Design	15%	
4. Prolog Implementation and Programming	15%	
5. Forward Chaining Implementation and Demonstration	10%	
6. Backward Chaining Implementation and Demonstration	10%	
7. Procedural and Non-Procedural Paradigm Analysis	10%	
8. Unification, Backtracking and Inference Accuracy	10%	
9. Testing, Results and System Demonstration	5%	
10. Documentation, GitHub Repository and Technical Presentation	5%	
Total	100%	

Signature of the Faculty

Total Marks Awarded

ANSWER ANY ONE OF THE FOLLOWING

Tool: SWI-Prolog

Usage of Prolog: Students shall use **Prolog** to develop a rule-based expert system by representing domain knowledge as **facts and production rules**, and implementing reasoning through **forward and backward chaining**. The project should demonstrate Prolog's **declarative/non-procedural nature**, logical inference, rule matching, and automated conclusion generation.

S.No.	Scenario-Based Expert System Development Question
1	Automobile Fault Diagnosis System: Design an expert system that identifies probable vehicle faults based on symptoms such as engine noise, starting problems, warning indicators, abnormal vibration, and reduced mileage. Implement the knowledge base using production rules and compare forward and backward reasoning.

1. Problem Analysis and Domain Understanding

Modern vehicles fail in ways that are diagnosable from a limited, observable set of symptoms: unusual engine noises, starting difficulties, dashboard warning indicators, abnormal vibration, and changes in fuel mileage or exhaust colour. A human mechanic reasons from these observations to a probable underlying fault using rules of thumb acquired through training and experience. This project builds a rule-based expert system that performs the same task automatically: given a set of observed symptoms, it identifies the most probable vehicle fault(s) and recommends a corrective action, and it can also be asked to verify or refute a specific hypothesised fault.

The problem domain was analysed along four dimensions, as required for a production-rule expert system:

- Entities: the vehicle (as an implicit single case), its symptoms, and its faults.
- Facts: atomic, observable statements supplied by the user/mechanic, e.g. `symptom(no_crank)`, `symptom(dim_headlights)`.
- Conditions / symptom categories: engine noise (knocking, squealing, grinding, ticking, hissing), starting problems (no crank, cranks-but-no-start, hard start, stalls after starting, slow crank, clicking on start), warning indicators (check-engine, battery, oil, temperature lights), abnormal vibration (idle, high-speed, braking, steering) and reduced mileage / smoke colour / smell observations.
- Goals / conclusions: 15 primitive vehicle faults (e.g. `weak_battery`, `faulty_alternator`, `worn_timing_belt`, `wheel_misalignment`) plus one second-level, derived conclusion (`electrical_system_failure`) that only follows once two lower-level faults have themselves been established — included specifically to demonstrate multi-step, chained inference rather than single-level rule matching.

Scope decisions: the system is intentionally symptom-driven and non-invasive (it does not require sensor/OBD-II data), mirroring an initial "first opinion" consultation a mechanic would give before physical inspection. Each conclusion also carries a plain-language recommended action, so the system's output is directly actionable.

2. Knowledge Acquisition and Knowledge Representation

Domain knowledge for automotive fault diagnosis is well documented in vehicle service manuals and standard automotive-technology references, which describe canonical symptom-to-cause associations (for example: a clicking sound on ignition combined with a failure to crank is a textbook indicator of a faulty starter motor, while dim headlights together with a slow crank and no crank at all points to a discharged battery rather than the starter itself). This project's knowledge base was compiled from such generally accepted automotive diagnostic heuristics, organised into IF (symptom conjunction) → THEN

(fault) associations — the standard acquisition method for a small expert system, sometimes called manual/interview-based elicitation when a domain expert is available, and document-based elicitation when service literature is used instead.

Two representation choices were made deliberately:

- Facts and rules are represented as ordinary Prolog clauses (symptom/1 facts, rule/1 production rules), giving each unit of knowledge an explicit, inspectable, declarative form rather than being buried in procedural code.
- A second predicate layer, recommend/2, associates each fault with a corrective action — separating "what is wrong" (diagnostic knowledge) from "what to do about it" (prescriptive knowledge), which keeps the knowledge base modular and easy to extend.

This yields a knowledge specification of 26 primitive symptom terms, 16 production rules (15 primitive + 1 derived/meta rule) and 16 recommendation facts — summarised in Section 3.

3. Production Rule / Knowledge Base Design

The knowledge base is organised as a flat set of independent IF–THEN production rules (Table 1), which is the representation format prescribed for this project. Rule R16 is marked with an asterisk because its condition is not a raw symptom but two previously derived fault/1 conclusions — it exists specifically to test and demonstrate multi-level chaining (a conclusion feeding another rule as a new premise), which a single-pass matcher would miss but the fixpoint forward-chaining engine in Section 5 correctly handles.

Rule	IF (symptom conditions)	THEN (conclusion)
R1	dim_headlights $\wedge$ slow_engine_crank $\wedge$ no_crank	fault(weak_battery)
R2	clicking_sound_on_start $\wedge$ no_crank	fault(faulty_starter_motor)
R3	battery_light $\wedge$ dim_headlights $\wedge$ stalls_after_start	fault(faulty_alternator)
R4	rough_idle $\wedge$ hard_start $\wedge$ check_engine_light	fault(worn_spark_plugs)
R5	reduced_mileage $\wedge$ black_smoke $\wedge$ rough_idle	fault(clogged_fuel_injector)
R6	cranks_no_start $\wedge$ fuel_smell	fault(faulty_fuel_pump)
R7	engine_ticking $\wedge$ check_engine_light $\wedge$ stalls_after_start	fault(worn_timing_belt)

R8	steering_vibration $\wedge$ high_speed_vibration	fault(wheel_misalignment)
R9	idle_vibration $\wedge$ engine_knocking	fault(worn_engine_mounts)
R10	reduced_mileage $\wedge$ check_engine_light $\wedge$ black_smoke	fault(faulty_oxygen_sensor)
R11	reduced_mileage $\wedge$ rough_idle $\wedge$ hard_start	fault(clogged_air_filter)
R12	braking_vibration $\wedge$ engine_squealing	fault(worn_brake_rotor)
R13	oil_light $\wedge$ low_oil_level	fault(low_engine_oil)
R14	temp_light $\wedge$ engine_overheating $\wedge$ white_smoke	fault(overheating_cooling_system)
R15	high_speed_vibration $\wedge$ engine_grinding	fault(worn_cv_joint)
R16*	fault(weak_battery) $\wedge$ fault(faulty_alternator)	fault(electrical_system_failure)

Table 1: Production rules in the automobile fault knowledge base.

Design properties verified during construction: (a) Consistency — no two rules conclude contradictory faults from an identical symptom set; (b) Non-redundancy — every rule contributes at least one condition not shared by all other rules for the same fault; (c) Completeness within scope — every one of the five symptom categories named in the assignment brief (engine noise, starting problems, warning indicators, abnormal vibration, reduced mileage) is covered by at least three rules; (d) Explainability — every rule maps to exactly one recommend/2 fact, so a conclusion is always paired with actionable advice.

4. Prolog Implementation and Programming

The system is implemented in SWI-Prolog (v9.0) as a single module, auto_diagnosis.pl. Three dynamic predicates hold the mutable state of a consultation:

```
:- dynamic(symptom/1). % observed symptom facts, asserted per case
:- dynamic(fault/1).   % derived conclusions, asserted by the inference engines
:- dynamic(known/2).  % memoised yes/no answers, used by backward chaining
```

A key design decision was to separate the production rules themselves (declared under rule/1) from the fact store of derived conclusions (fault/1, populated only by assert). Writing the rules directly as fault/1 clauses was tried first and rejected: because fault(X) would then succeed by direct resolution the moment its symptoms were known, the forward-chaining fixpoint test "\+ fault(F)" (has F not yet been

derived?) could never distinguish 'derivable' from 'already recorded', so no new fact was ever asserted. Splitting the predicates — rule/1 for "is this provable right now" and fault/1 for "has this been recorded as derived" — fixed this and is what allows rule R16 to chain correctly off two independently-derived fault/1 facts.

Representative rule and recommendation clauses:

```
rule(weak_battery) :-
    symptom(dim_headlights),
    symptom(slow_engine_crank),
    symptom(no_crank).

rule(electrical_system_failure) :-
    fault(weak_battery),
    fault(faulty_alternator).

recommend(weak_battery, 'Charge or replace the battery; inspect terminals for corrosion.').
```

Every symptom (symptom/1), fault (fault/1) and rule (rule/1) is a standard first-order Prolog term, so the engines rely entirely on Prolog's built-in unification and SLD-resolution to match conditions against known facts — no external rule-matching library is used. The complete, runnable source listing is provided in Appendix A.

5. Forward Chaining Implementation and Demonstration

Forward chaining is data-driven: reasoning starts from whatever facts are known and repeatedly fires any rule whose body is fully satisfied, asserting its conclusion as a new fact, until a full pass over the rule set produces nothing new (a fixpoint). This is implemented by forward_chain_step/1, which scans every rule name in fault_rule_list/1, collects every fault F such that F is not yet asserted but rule(F) now succeeds, asserts each one (printing an explanation trace), and recurses — so that newly asserted facts (such as weak_battery and faulty_alternator) can in turn satisfy a later rule (R16) within the same consultation:

```
forward_chain_step(_):-
    fault_rule_list(Faults),
    findall(F, ( member(F, Faults), \+ fault(F), call(rule, F) ), NewFaults),
    ( NewFaults == []
    -> true                % fixpoint reached
    ; ( forall(member(F, NewFaults), (assertz(fault(F)), assertz(fired(F)))),
        forward_chain_step(_) ) % re-scan for newly-enabled rules
    ).
```

run_forward/1 asserts a list of observed symptoms, invokes forward_chain/0, and prints every derived

fault with its recommended action. Below is the actual console output produced by running the compiled program (captured via `swipl -q -f auto_diagnosis.pl`):

```
PS C:\Users\arathi\Downloads> & "C:\Program Files\swipl\bin\swipl.exe" -q -f auto_diagnosis.pl

===== FORWARD CHAINING RUN =====
Asserting observed symptoms: [dim_headlights,slow_engine_crank,no_crank]
[Forward-Chain] Rule fired -> fault(weak_battery) derived

Concluded faults:
- weak_battery => Charge or replace the battery; inspect terminals for corrosion.

===== FORWARD CHAINING RUN =====
Asserting observed symptoms: [reduced_mileage,black_smoke,rough_idle,check_engine_light,hard_start]
[Forward-Chain] Rule fired -> fault(worn_spark_plugs) derived
[Forward-Chain] Rule fired -> fault(clogged_fuel_injector) derived
[Forward-Chain] Rule fired -> fault(faulty_oxygen_sensor) derived
[Forward-Chain] Rule fired -> fault(clogged_air_filter) derived

Concluded faults:
- worn_spark_plugs => Replace spark plugs and inspect ignition coils.
- clogged_fuel_injector => Clean or replace fuel injectors; use injector cleaner additive.
- faulty_oxygen_sensor => Replace the oxygen (O2) sensor and clear the ECU fault code.
- clogged_air_filter => Replace the air filter.

===== FORWARD CHAINING RUN =====
Asserting observed symptoms: [battery_light,dim_headlights,stalls_after_start,slow_engine_crank,no_crank]
[Forward-Chain] Rule fired -> fault(weak_battery) derived
[Forward-Chain] Rule fired -> fault(faulty_alternator) derived
[Forward-Chain] Rule fired -> fault(electrical_system_failure) derived

Concluded faults:
- weak_battery => Charge or replace the battery; inspect terminals for corrosion.
- faulty_alternator => Test alternator output; replace if charging voltage is low.
- electrical_system_failure => Have the entire charging/electrical system tested by a technician.

===== FORWARD CHAINING RUN =====
Asserting observed symptoms: [engine_hissing]
No fault could be concluded from the given symptoms.
```

```
===== FORWARD CHAINING RUN =====

Asserting observed symptoms: [dim_headlights,slow_engine_crank,no_crank]

[Forward-Chain] Rule fired -> fault(weak_battery) derived


Concluded faults:

- weak_battery => Charge or replace the battery; inspect terminals for corrosion.


===== FORWARD CHAINING RUN =====

Asserting observed symptoms: [reduced_mileage,black_smoke,rough_idle,check_engine_light,hard_start]

[Forward-Chain] Rule fired -> fault(worn_spark_plugs) derived
[Forward-Chain] Rule fired -> fault(clogged_fuel_injector) derived
[Forward-Chain] Rule fired -> fault(faulty_oxygen_sensor) derived
[Forward-Chain] Rule fired -> fault(clogged_air_filter) derived


Concluded faults:

- worn_spark_plugs => Replace spark plugs and inspect ignition coils.
- clogged_fuel_injector => Clean or replace fuel injectors; use injector cleaner additive.
- faulty_oxygen_sensor => Replace the oxygen (O2) sensor and clear the ECU fault code.
- clogged_air_filter => Replace the air filter.
```

```

===== FORWARD CHAINING RUN =====
Asserting observed symptoms: [battery_light,dim_headlights,stalls_after_start,slow_engine_crank,no_crank]
[Forward-Chain] Rule fired -> fault(weak_battery) derived
[Forward-Chain] Rule fired -> fault(faulty_alternator) derived
[Forward-Chain] Rule fired -> fault(electrical_system_failure) derived

Concluded faults:
- weak_battery => Charge or replace the battery; inspect terminals for corrosion.
- faulty_alternator => Test alternator output; replace if charging voltage is low.
- electrical_system_failure => Have the entire charging/electrical system tested by a technician.

===== FORWARD CHAINING RUN =====
Asserting observed symptoms: [engine_hissing]
No fault could be concluded from the given symptoms.

```

The third run demonstrates chained, multi-level forward inference exactly as intended: from five raw symptoms the engine first derives `weak_battery` and `faulty_alternator` directly from symptoms, and only then — on its second pass over the rule set — derives `electrical_system_failure` from those two freshly-asserted facts. The fourth run (a single, insufficient symptom) correctly reports that no conclusion can be drawn, confirming the engine does not over-generalise from partial evidence.

6. Backward Chaining Implementation and Demonstration

Backward chaining is goal-driven: reasoning starts from a hypothesised conclusion (a specific fault) and works backwards through the rules to the symptoms that would justify it, asking about any symptom whose truth is not yet known. This is implemented as an explicit meta-interpreter, `solve/1`, rather than relying on Prolog's default top-level resolution directly, so that the goal-driven expansion and the "ask only if unknown" behaviour are visible and explainable:

```

solve(fault(F)) :-
    !,
    clause(rule(F), Body),
    solve(Body),
    ( fired(F) -> true ; assertz(fired(F)) ),
    ( fault(F) -> true ; assertz(fault(F)) ). % memoise for chained goals (R16)
solve((A, B)) :- !, solve(A), solve(B).
solve(symptom(S)) :- !, ask(S).

ask(Symptom) :- known(Symptom, yes), !.
ask(Symptom) :- known(Symptom, no), !, fail.
ask(Symptom) :- ... % otherwise query the user (or a preset answer list) and memoise it

```

`known/2` memoises every answer the moment it is given, so the same symptom is never asked twice

even if it appears in more than one rule reached during the search. For repeatable, automated testing (used throughout this report) ask/1 consults a preset answer list instead of the keyboard; the identical predicate also supports a live, interactive yes/no consultation via consult_interactive/0. Below is the actual console output from five backward-chaining runs, including one designed to fail (a refutation) and one that proves the chained meta-rule R16 by recursing through two nested fault(...) subgoals:

```
===== BACKWARD CHAINING RUN: goal = fault(worn_brake_rotor) =====
[Backward-Chain] Q: is 'braking_vibration' present? (auto) -> yes
[Backward-Chain] Q: is 'engine_squealing' present? (auto) -> yes

GOAL PROVED: fault(worn_brake_rotor) confirmed.
Recommendation: Inspect and machine/replace brake rotors and pads.

===== BACKWARD CHAINING RUN: goal = fault(faulty_starter_motor) =====
[Backward-Chain] Q: is 'clicking_sound_on_start' present? (auto) -> yes
```

```
[Backward-Chain] Q: is 'no_crank' present? (auto) -> no

GOAL NOT PROVED: insufficient evidence for fault(faulty_starter_motor).

===== BACKWARD CHAINING RUN: goal = fault(wheel_misalignment) =====
[Backward-Chain] Q: is 'steering_vibration' present? (auto) -> yes
[Backward-Chain] Q: is 'high_speed_vibration' present? (auto) -> yes

GOAL PROVED: fault(wheel_misalignment) confirmed.
Recommendation: Perform a wheel alignment and check tyre wear.
```

```
===== BACKWARD CHAINING RUN: goal = fault(worn_timing_belt) =====
[Backward-Chain] Q: is 'engine_ticking' present? (auto) -> yes
[Backward-Chain] Q: is 'check_engine_light' present? (auto) -> yes
[Backward-Chain] Q: is 'stalls_after_start' present? (auto) -> yes

GOAL PROVED: fault(worn_timing_belt) confirmed.
Recommendation: Inspect and replace the timing belt immediately - risk of engine damage.

===== BACKWARD CHAINING RUN: goal = fault(electrical_system_failure) =====
[Backward-Chain] Q: is 'dim_headlights' present? (auto) -> yes
[Backward-Chain] Q: is 'slow_engine_crank' present? (auto) -> yes
[Backward-Chain] Q: is 'no_crank' present? (auto) -> yes
[Backward-Chain] Q: is 'battery_light' present? (auto) -> yes
[Backward-Chain] Q: is 'stalls_after_start' present? (auto) -> yes

GOAL PROVED: fault(electrical_system_failure) confirmed.
Recommendation: Have the entire charging/electrical system tested by a technician.

===== ALL TESTS COMPLETE =====
Warning: [Thread main] c:/users/arthi/downloads/auto_diagnosis.pl:367: main
Warning: [Thread main] Initialization goal called halt(0).
Warning: [Thread main] The program entry point should be called using initialization/2.
Warning: [Thread main] Consider using library(main).
```

```
===== BACKWARD CHAINING RUN: goal = fault(worn_brake_rotor) =====
[Backward-Chain] Q: is 'braking_vibration' present? (auto) -> yes
[Backward-Chain] Q: is 'engine_squealing' present? (auto) -> yes

GOAL PROVED: fault(worn_brake_rotor) confirmed.
```

Recommendation: Inspect and machine/replace brake rotors and pads.

===== BACKWARD CHAINING RUN: goal = fault(faulty_starter_motor) =====

[Backward-Chain] Q: is 'clicking_sound_on_start' present? (auto) -> yes

[Backward-Chain] Q: is 'no_crank' present? (auto) -> no

GOAL NOT PROVED: insufficient evidence for fault(faulty_starter_motor).

===== BACKWARD CHAINING RUN: goal = fault(wheel_misalignment) =====

[Backward-Chain] Q: is 'steering_vibration' present? (auto) -> yes

[Backward-Chain] Q: is 'high_speed_vibration' present? (auto) -> yes

GOAL PROVED: fault(wheel_misalignment) confirmed.

Recommendation: Perform a wheel alignment and check tyre wear.

===== BACKWARD CHAINING RUN: goal = fault(worn_timing_belt) =====

[Backward-Chain] Q: is 'engine_ticking' present? (auto) -> yes

[Backward-Chain] Q: is 'check_engine_light' present? (auto) -> yes

[Backward-Chain] Q: is 'stalls_after_start' present? (auto) -> yes

GOAL PROVED: fault(worn_timing_belt) confirmed.

Recommendation: Inspect and replace the timing belt immediately - risk of engine damage.

===== BACKWARD CHAINING RUN: goal = fault(electrical_system_failure) =====

[Backward-Chain] Q: is 'dim_headlights' present? (auto) -> yes

[Backward-Chain] Q: is 'slow_engine_crank' present? (auto) -> yes

[Backward-Chain] Q: is 'no_crank' present? (auto) -> yes

[Backward-Chain] Q: is 'battery_light' present? (auto) -> yes

[Backward-Chain] Q: is 'stalls_after_start' present? (auto) -> yes

GOAL PROVED: fault(electrical_system_failure) confirmed.

Recommendation: Have the entire charging/electrical system tested by a technician.

===== ALL TESTS COMPLETE =====

The second run (faulty_starter_motor) is a deliberate refutation test: the user confirms clicking_sound_on_start but denies no_crank, so rule R2's conjunction cannot be satisfied and the engine correctly reports the goal as not proved, rather than guessing. The fifth run proves electrical_system_failure purely by backward reasoning — solve/1 expands rule(electrical_system_failure), meets the subgoal fault(weak_battery), and correctly recurses into solve(fault(weak_battery)) to prove it from raw symptoms before continuing to fault(faulty_alternator), demonstrating goal-driven reasoning across two rule levels.

7. Procedural and Non-Procedural Paradigm Analysis

This project's knowledge base could equally be expressed as a procedural program — for example, a chain of nested if-statements in Python that tests symptom flags and prints a diagnosis. Table 2 contrasts that hypothetical procedural implementation with the declarative Prolog implementation actually used.

Aspect	Procedural approach	Non-procedural (declarative/Prolog) approach
Control flow	Programmer explicitly sequences every IF/ELSE/loop that walks the decision tree (e.g. nested if-statements in Python/Java).	Programmer states only WHAT relationships hold (rule/1 facts); Prolog's SLD-resolution engine decides HOW and in what order to try them.
Adding new knowledge	Requires locating and editing the control-flow code, risking side effects on unrelated branches.	Requires only adding one new rule/1 or symptom/1 clause; existing clauses and the resolution engine are untouched.
Reasoning direction	Usually hard-codes one direction (e.g. only symptom→fault); reversing to "which symptoms would prove fault X" needs new code.	The same rule/1 clauses are reused unmodified for BOTH forward chaining (this project's forward_chain_step/1) and backward chaining (solve/1) — direction is a property of the engine, not the knowledge.
Search / backtracking	Must be manually implemented (explicit stacks/recursion) if multiple solution paths exist.	Backtracking is built into the language; solve/1 and clause/2 automatically try alternative rules for the same fault without extra code.
Matching conditions	Explicit equality checks and boolean logic written by hand for every condition.	Unification, Prolog's built-in pattern-matching, does this automatically when a goal such as symptom(no_crank)

		is matched against the fact base.
--	--	-----------------------------------

Table 2: Procedural vs. non-procedural paradigm comparison.

The clearest practical demonstration of the declarative advantage in this codebase is that a single set of rule/1 clauses drives two entirely different reasoning strategies (`forward_chain_step/1` and `solve/1`) without any duplication of domain knowledge — something that would require two separately-coded and separately-maintained procedures under a procedural paradigm.

8. Unification, Backtracking and Inference Accuracy

Unification: every call such as `symptom(no_crank)` or `clause(rule(F), Body)` is resolved by Prolog attempting to unify the call with the head of a stored clause. Because `F` is a bound variable at the point `solve(fault(F))` is invoked (it comes from the user's chosen goal, e.g. `fault(worn_timing_belt)`), `clause(rule(F), Body)` unifies `F` against each stored `rule(...)` head in turn until an exact match binds `Body` to that rule's specific condition conjunction — this is how the same `solve/1` clause works for all sixteen different rules without being rewritten for each one.

Backtracking: `findall/3` inside `forward_chain_step/1` exhaustively backtracks over every member of `fault_rule_list/1`, retrying `rule(F)` for each candidate fault name and collecting every one whose body succeeds, rather than committing to the first match — necessary because more than one rule can legitimately fire from the same symptom set (as seen in the second forward-chaining test run in Section 5, where four independent faults were derived from one overlapping symptom set). In `solve/1`, `clause/2`'s own backtracking would similarly allow an alternative rule for the same fault to be tried automatically if the first rule's body failed — the cut (!) immediately after the head match is used deliberately in `solve/1` to commit to structural cases (`fault/1` vs. `conjunction` vs. `symptom/1`) without cutting away this rule-alternative backtracking, since each fault currently has only one production rule.

Inference accuracy was validated against seven test scenarios (Section 9): true-positive derivations (a fault correctly concluded when all its conditions hold), true-negative refutations (a fault correctly reported as not provable when one condition is denied), a no-conclusion case (symptoms too sparse to match any rule), and a chained derivation (R16) proved consistently by both engines. All seven produced the expected result, with zero incorrect conclusions and zero missed derivations.

9. Testing, Results and System Demonstration

Seven scenarios were run through the automated test harness (`run_all_tests/0`), covering both engines, positive derivations, a refutation, a no-conclusion case, and chained (multi-level) inference. Table 3 summarises the results; the full, unedited console transcript is reproduced in Sections 5 and 6 and in Appendix B.

#	Engine	Input (symptoms / goal)	Expected result	Actual result
1	Forward	dim_headlights, slow_engine_crank, no_crank	weak_battery	weak_battery — PASS
2	Forward	reduced_mileage, black_smoke, rough_idle, check_engine_light, hard_start	4 faults (spark plugs, injector, O2 sensor, air filter)	4 faults derived — PASS
3	Forward	battery_light, dim_headlights, stalls_after_start, slow_engine_crank, no_crank	weak_battery + faulty_alternator + chained electrical_system_failure	all 3 derived, R16 fired 2nd pass — PASS
4	Forward	engine_hissing (only)	no conclusion	no fault concluded — PASS
5	Backward	goal=worn_brake_rotor; braking_vibration=Y, engine_squealing=Y	goal proved	GOAL PROVED — PASS
6	Backward	goal=faulty_starter_motor; clicking_sound_on_start=Y, no_crank=N	goal NOT proved (refutation)	GOAL NOT PROVED — PASS
7	Backward	goal=electrical_system_failure; 5 supporting symptoms=Y	goal proved via chained R16	GOAL PROVED — PASS

Table 3: Test case summary — all 7/7 scenarios passed.

The tests were executed with SWI-Prolog 9.0.4 using the command "swipl -q -f auto_diagnosis.pl", which loads the knowledge base and automatically runs the full test suite via an initialization/1 directive — making the demonstration fully repeatable without any manual keyboard input, which is also how the transcripts in this report were produced. For a live, interactive viva demonstration, consult_interactive/0 can instead be called after loading the file, which asks about each symptom in turn and then runs forward chaining on the answers given.

10. Documentation, GitHub Repository and Technical Presentation

Recommended repository structure for submission:

```
automobile-fault-diagnosis-expert-system/  
├── auto_diagnosis.pl    (complete knowledge base + both inference engines)  
├── README.md           (problem statement, how to run, sample output)  
├── report/  
│   └── Automobile_Fault_Diagnosis_Report.docx  
└── LICENSE
```

Suggested steps to publish and present the project:

- git init, then git add auto_diagnosis.pl README.md, and an initial commit with a clear message.
- Create the GitHub repository, add it as the origin remote, and push: git remote add origin <repo-url> then git push -u origin main.
- README.md should state the problem, list the symptom vocabulary and the 16 production rules, give the exact run command (swipl -q -f auto_diagnosis.pl), and paste a short sample of the console output.
- During the technical presentation/viva: (1) explain the domain and rule design (Sections 1–3), (2) show the rule/1 vs fault/1 separation and why it was necessary (Section 4), (3) live-run forward chaining on a fresh symptom set, (4) live-run backward chaining with consult_interactive/0 or run_backward/2 on a hypothesis the examiner names, and (5) walk through one refutation case to show the system does not over-claim.

11. Conclusion

The completed system satisfies every stage of the required report format: a defined problem domain, an acquired and represented knowledge base of 16 production rules, a fully working SWI-Prolog implementation, two independent and correctly-demonstrated inference engines (data-driven forward chaining and goal-driven backward chaining, including a chained/meta rule proved by both), an explicit paradigm comparison, and a seven-case test suite that passed in full. The forward/backward split of rule/1 vs. fault/1 — discovered and fixed during development — is a good illustration of how Prolog's unification-driven resolution, while powerful, still requires careful predicate design when a knowledge base is meant to support more than one reasoning strategy over the same rules.

Appendix A: Complete Prolog Source Listing

(auto_diagnosis.pl)

```
/* =====
AUTOMOBILE FAULT DIAGNOSIS EXPERT SYSTEM
Language : SWI-Prolog
Purpose  : Identify probable vehicle faults from symptoms
           (engine noise, starting problems, warning
           indicators, abnormal vibration, reduced mileage)
           using production rules, with both FORWARD
           CHAINING (data-driven) and BACKWARD CHAINING
           (goal-driven) inference engines.
===== */

:- dynamic(symptom/1).      % observed symptom facts
:- dynamic(fault/1).        % derived fault conclusions
:- dynamic(known/2).        % memoised answers for backward chaining: known(Symptom,yes/no)
:- dynamic(fired/1).        % rules that fired, for explanation trace

/* -----
STAGE 1-2 : DOMAIN MODEL / SYMPTOM VOCABULARY
-----

engine_knocking, engine_squealing, engine_grinding,
engine_ticking, engine_hissing      -- engine noise
no_crank, cranks_no_start, hard_start,
stalls_after_start, slow_engine_crank,
clicking_sound_on_start              -- starting problems
check_engine_light, battery_light,
oil_light, temp_light                -- warning indicators
idle_vibration, high_speed_vibration,
braking_vibration, steering_vibration -- abnormal vibration
reduced_mileage, black_smoke, white_smoke,
fuel_smell, rough_idle, dim_headlights,
low_oil_level, engine_overheating    -- other observations
----- */

/* =====
STAGE 3 : KNOWLEDGE BASE - PRODUCTION RULES (IF-THEN)
Each fault rule is written as:
    fault(Name) :- Condition1, Condition2, ...
Conditions are symptom/1 facts OR other fault/1 conclusions,
which allows multi-level (chained) inference.
===== */

% NOTE: production rules are written under rule/1. fault/1 stays a
% pure dynamic FACT store, holding only conclusions that an
```

```
% inference engine has actually derived/asserted. Keeping the two
% separate lets forward chaining tell "provable-by-rule" apart
% from "already-derived-and-recorded" (needed for its fixpoint
% termination test, \+ fault(F)).
```

```
% R1: Weak / discharged battery
```

```
rule(weak_battery) :-
    symptom(dim_headlights),
    symptom(slow_engine_crank),
    symptom(no_crank).
```

```
% R2: Faulty starter motor
```

```
rule(faulty_starter_motor) :-
    symptom(clicking_sound_on_start),
    symptom(no_crank).
```

```
% R3: Faulty alternator
```

```
rule(faulty_alternator) :-
    symptom(battery_light),
    symptom(dim_headlights),
    symptom(stalls_after_start).
```

```
% R4: Worn spark plugs
```

```
rule(worn_spark_plugs) :-
    symptom(rough_idle),
    symptom(hard_start),
    symptom(check_engine_light).
```

```
% R5: Clogged fuel injector
```

```
rule(clogged_fuel_injector) :-
    symptom(reduced_mileage),
    symptom(black_smoke),
    symptom(rough_idle).
```

```
% R6: Faulty fuel pump
```

```
rule(faulty_fuel_pump) :-
    symptom(cranks_no_start),
    symptom(fuel_smell).
```

```
% R7: Worn timing belt
```

```
rule(worn_timing_belt) :-
    symptom(engine_ticking),
    symptom(check_engine_light),
    symptom(stalls_after_start).
```

```
% R8: Wheel misalignment
```



```
rule(wheel_misalignment) :-  
    symptom(steering_vibration),  
    symptom(high_speed_vibration).
```

% R9: Worn engine mounts

```
rule(worn_engine_mounts) :-  
    symptom(idle_vibration),  
    symptom(engine_knocking).
```

% R10: Faulty oxygen sensor

```
rule(faulty_oxygen_sensor) :-  
    symptom(reduced_mileage),  
    symptom(check_engine_light),  
    symptom(black_smoke).
```

% R11: Clogged air filter

```
rule(clogged_air_filter) :-  
    symptom(reduced_mileage),  
    symptom(rough_idle),  
    symptom(hard_start).
```

% R12: Worn brake rotor

```
rule(worn_brake_rotor) :-  
    symptom(braking_vibration),  
    symptom(engine_squealing).
```

% R13: Low engine oil

```
rule(low_engine_oil) :-  
    symptom(oil_light),  
    symptom(low_oil_level).
```

% R14: Overheating cooling system

```
rule(overheating_cooling_system) :-  
    symptom(temp_light),  
    symptom(engine_overheating),  
    symptom(white_smoke).
```

% R15: Worn CV joint

```
rule(worn_cv_joint) :-  
    symptom(high_speed_vibration),  
    symptom(engine_grinding).
```

% ---- Second-level (meta) production rule: demonstrates

% ---- chaining over DERIVED facts, not just raw symptoms.

% R16: Electrical system failure (derived from two other faults)

```
rule(electrical_system_failure) :-
```

```

fault(weak_battery),
fault(faulty_alternator).

/* =====

STAGE 4a : ADVICE / ACTION PRODUCTION RULES
IF fault(X) THEN recommend(Action)
===== */

recommend(weak_battery,      'Charge or replace the battery; inspect terminals for corrosion.').
recommend(faulty_starter_motor, 'Inspect/replace the starter motor and check its wiring.').
recommend(faulty_alternator,   'Test alternator output; replace if charging voltage is low.').
recommend(worn_spark_plugs,     'Replace spark plugs and inspect ignition coils.').
recommend(clogged_fuel_injector, 'Clean or replace fuel injectors; use injector cleaner additive.').
recommend(faulty_fuel_pump,     'Test fuel pump pressure; replace pump if pressure is low.').
recommend(worn_timing_belt,     'Inspect and replace the timing belt immediately - risk of engine damage.').
recommend(wheel_misalignment,  'Perform a wheel alignment and check tyre wear.').
recommend(worn_engine_mounts,  'Inspect and replace worn engine/transmission mounts.').
recommend(faulty_oxygen_sensor, 'Replace the oxygen (O2) sensor and clear the ECU fault code.').
recommend(clogged_air_filter,  'Replace the air filter.').
recommend(worn_brake_rotor,    'Inspect and machine/replace brake rotors and pads.').
recommend(low_engine_oil,      'Top up engine oil immediately and check for leaks.').
recommend(overheating_cooling_system, 'Check coolant level, radiator, water pump and thermostat urgently.').
recommend(worn_cv_joint,       'Inspect and replace the CV joint / drive-shaft boot.').
recommend(electrical_system_failure, 'Have the entire charging/electrical system tested by a technician.').

/* =====

STAGE 4b : FORWARD CHAINING ENGINE (data-driven, non-recursive
descent - iterates the rule set repeatedly over KNOWN facts
until no new fault/1 can be derived: classic fixpoint forward
chaining used in production-rule systems).
===== */

fault_rule_list([
    weak_battery, faulty_starter_motor, faulty_alternator,
    worn_spark_plugs, clogged_fuel_injector, faulty_fuel_pump,
    worn_timing_belt, wheel_misalignment, worn_engine_mounts,
    faulty_oxygen_sensor, clogged_air_filter, worn_brake_rotor,
    low_engine_oil, overheating_cooling_system, worn_cv_joint,
    electrical_system_failure
]).

forward_chain :-
    retractall(fired(_)),
    forward_chain_step(_),
    !.

% Repeatedly scan every rule; assert any newly provable fault;
% stop when a full pass adds nothing new (fixpoint reached).

```

```

forward_chain_step(_):-
    fault_rule_list(Faults),
    findall(F,
        ( member(F, Faults),
          \+ fault(F),      % not yet recorded as derived
          call(rule, F)     % but its production rule now fires
        ),
        NewFaults),
    ( NewFaults == []
    -> true                % fixpoint reached, stop
    ; ( forall(member(F, NewFaults),
        ( assertz(fault(F)),
          assertz(fired(F)),
          format(" [Forward-Chain] Rule fired -> fault(~w) derived~n", [F])
        )
      ),
      forward_chain_step(_)    % re-scan: new facts may enable more rules (e.g. R16)
    )
    ).

```

```

run_forward(Symptoms):-
    format("~n===== FORWARD CHAINING RUN =====~n"),
    format("Asserting observed symptoms: ~w~n", [Symptoms]),
    retractall(symptom(_)), retractall(fault(_)),
    forall(member(S, Symptoms), assertz(symptom(S))),
    forward_chain,
    findall(F, fault(F), Faults),
    ( Faults == []
    -> format("No fault could be concluded from the given symptoms.~n")
    ; ( format("~nConcluded faults:~n"),
        forall(member(F, Faults),
            ( recommend(F, Advice)
            -> format(" - ~w => ~w~n", [F, Advice])
            ; format(" - ~w~n", [F])
            ))
        )
    ).

```

```

/* =====
STAGE 4c : BACKWARD CHAINING ENGINE (goal-driven).
Tries to prove a specific fault by working BACKWARDS from the
goal through the rules to the symptoms. Any symptom whose
truth is not yet known is treated as an ASKABLE leaf: the
engine queries the user (or a supplied answer list, for
automated / repeatable testing) and memoises the answer in
known/2 so it is never asked twice.
===== */

```

```

:- dynamic(answers/1). % preset answers used during automated testing

% ask/1 : resolve the truth of a single symptom
ask(Symptom) :-
    known(Symptom, yes), !.
ask(Symptom) :-
    known(Symptom, no), !, fail.
ask(Symptom) :-
    answers(List), !, % automated / non-interactive mode
    ( member(Symptom-Answer, List)
    -> true
    ; Answer = no
    ),
    assertz(known(Symptom, Answer)),
    format(" [Backward-Chain] Q: is '~w' present? (auto) -> ~w~n", [Symptom, Answer]),
    Answer == yes.
ask(Symptom) :-
    format(" [Backward-Chain] Is '~w' present? (yes/no): ", [Symptom]),
    read(Answer),
    assertz(known(Symptom, Answer)),
    Answer == yes.

% solve/1 : explicit backward-chaining meta-interpreter.
% Goal-driven: starts at fault(Goal) and expands it using the
% SAME production rules declared above, recursing into subgoals
% (which may themselves be faults, e.g. electrical_system_failure)
% and only consulting ask/1 for primitive symptom leaves.
solve(fault(F)) :-
    !,
    clause(rule(F), Body),
    solve(Body),
    ( fired(F) -> true ; assertz(fired(F)) ),
    ( fault(F) -> true ; assertz(fault(F)) ). % memoise so R16-style chained goals reuse it
solve((A, B)) :-
    !,
    solve(A),
    solve(B).
solve(symptom(S)) :-
    !,
    ask(S).

run_backward(Fault, AnswerList) :-
    format("~n===== BACKWARD CHAINING RUN: goal = fault(~w) =====~n", [Fault]),
    retractall(known(_, _)), retractall(fired(_)), retractall(fault(_)), retractall(answers(_)),
    assertz(answers(AnswerList)),
    ( solve(fault(Fault))

```

```

-> ( recommend(Fault, Advice)
    -> format("~nGOAL PROVED: fault(~w) confirmed.~nRecommendation: ~w~n", [Fault, Advice])
    ; format("~nGOAL PROVED: fault(~w) confirmed.~n", [Fault])
    )
; format("~nGOAL NOT PROVED: insufficient evidence for fault(~w).~n", [Fault])
).

/* =====
STAGE 7 : TEST HARNESS
Demonstrates forward chaining on complete symptom sets and
backward chaining proving/refuting specific hypotheses.
===== */

test_forward_1 :-
    run_forward([dim_headlights, slow_engine_crank, no_crank]).

test_forward_2 :-
    run_forward([reduced_mileage, black_smoke, rough_idle,
        check_engine_light, hard_start]).

test_forward_3 :-
    run_forward([battery_light, dim_headlights, stalls_after_start,
        slow_engine_crank, no_crank]).

test_forward_4 :-
    run_forward([engine_hissing]). % no rule concludes -> tests "no conclusion" path

test_backward_1 :-
    run_backward(worn_brake_rotor,
        [braking_vibration=yes, engine_squealing=yes]).

test_backward_2 :-
    run_backward(faulty_starter_motor,
        [clicking_sound_on_start=yes, no_crank=no]). % expected to FAIL (refutation)

test_backward_3 :-
    run_backward(wheel_misalignment,
        [steering_vibration=yes, high_speed_vibration=yes]).

test_backward_4 :-
    % Proves worn_timing_belt with all evidence present -> should SUCCEED
    run_backward(worn_timing_belt,
        [engine_ticking=yes, check_engine_light=yes, stalls_after_start=yes]).

test_backward_5 :-

```

```

% Proves the CHAINED meta-rule R16: backward chaining must recurse
% through fault(weak_battery) and fault(faulty_alternator) subgoals
run_backward(electrical_system_failure,
    [dim_headlights=yes, slow_engine_crank=yes, no_crank=yes,
    battery_light=yes, stalls_after_start=yes]).

run_all_tests :-
    test_forward_1, test_forward_2, test_forward_3, test_forward_4,
    test_backward_1, test_backward_2, test_backward_3,
    test_backward_4, test_backward_5,
    format("~n===== ALL TESTS COMPLETE =====~n").

/* =====
STAGE 6 : SIMPLE INTERACTIVE CONSULTATION (user interface)
===== */

symptom_menu([
    dim_headlights, slow_engine_crank, no_crank, clicking_sound_on_start,
    battery_light, stalls_after_start, rough_idle, hard_start,
    check_engine_light, reduced_mileage, black_smoke, cranks_no_start,
    fuel_smell, engine_ticking, steering_vibration, high_speed_vibration,
    idle_vibration, engine_knocking, braking_vibration, engine_squealing,
    oil_light, low_oil_level, temp_light, engine_overheating,
    white_smoke, engine_grinding
]).

% consult_interactive/0 : asks about every symptom once, then runs
% forward chaining over the answers given (menu-driven demo mode).
consult_interactive :-
    format("~n--- Vehicle Fault Consultation (answer y/n) ---~n"),
    symptom_menu(Menu),
    findall(S,
        ( member(S, Menu),
          format("Is '~w' present? (y/n): ", [S]),
          read(Ans),
          ( Ans == y ; Ans == yes )
        ),
        Present),
    run_forward(Present).

/* Loading this file automatically runs the full automated test
suite (repeatable, no keyboard input needed - ideal for
grading/CI). For a LIVE menu-driven demonstration instead,
comment out the two lines below and, after loading the file
at the swipl prompt, call: ?- consult_interactive.    */
:- initialization(main).
main :-

```

```
run_all_tests,  
halt.
```

Appendix B: Full Console Transcript (all 7 test runs)

Captured verbatim from: swipl -q -f auto_diagnosis.pl

===== FORWARD CHAINING RUN =====

Asserting observed symptoms: [dim_headlights,slow_engine_crank,no_crank]

[Forward-Chain] Rule fired -> fault(weak_battery) derived

Concluded faults:

- weak_battery => Charge or replace the battery; inspect terminals for corrosion.

===== FORWARD CHAINING RUN =====

Asserting observed symptoms: [reduced_mileage,black_smoke,rough_idle,check_engine_light,hard_start]

[Forward-Chain] Rule fired -> fault(worn_spark_plugs) derived

[Forward-Chain] Rule fired -> fault(clogged_fuel_injector) derived

[Forward-Chain] Rule fired -> fault(faulty_oxygen_sensor) derived

[Forward-Chain] Rule fired -> fault(clogged_air_filter) derived

Concluded faults:

- worn_spark_plugs => Replace spark plugs and inspect ignition coils.

- clogged_fuel_injector => Clean or replace fuel injectors; use injector cleaner additive.

- faulty_oxygen_sensor => Replace the oxygen (O2) sensor and clear the ECU fault code.

- clogged_air_filter => Replace the air filter.

===== FORWARD CHAINING RUN =====

Asserting observed symptoms: [battery_light,dim_headlights,stalls_after_start,slow_engine_crank,no_crank]

[Forward-Chain] Rule fired -> fault(weak_battery) derived

[Forward-Chain] Rule fired -> fault(faulty_alternator) derived

[Forward-Chain] Rule fired -> fault(electrical_system_failure) derived

Concluded faults:

- weak_battery => Charge or replace the battery; inspect terminals for corrosion.

- faulty_alternator => Test alternator output; replace if charging voltage is low.

- electrical_system_failure => Have the entire charging/electrical system tested by a technician.

===== FORWARD CHAINING RUN =====

Asserting observed symptoms: [engine_hissing]

No fault could be concluded from the given symptoms.

===== BACKWARD CHAINING RUN: goal = fault(worn_brake_rotor) =====

[Backward-Chain] Q: is 'braking_vibration' present? (auto) -> yes

[Backward-Chain] Q: is 'engine_squealing' present? (auto) -> yes

GOAL PROVED: fault(worn_brake_rotor) confirmed.

Recommendation: Inspect and machine/replace brake rotors and pads.

===== BACKWARD CHAINING RUN: goal = fault(faulty_starter_motor) =====

[Backward-Chain] Q: is 'clicking_sound_on_start' present? (auto) -> yes

[Backward-Chain] Q: is 'no_crank' present? (auto) -> no

GOAL NOT PROVED: insufficient evidence for fault(faulty_starter_motor).

===== BACKWARD CHAINING RUN: goal = fault(wheel_misalignment) =====

[Backward-Chain] Q: is 'steering_vibration' present? (auto) -> yes

[Backward-Chain] Q: is 'high_speed_vibration' present? (auto) -> yes

GOAL PROVED: fault(wheel_misalignment) confirmed.

Recommendation: Perform a wheel alignment and check tyre wear.

===== BACKWARD CHAINING RUN: goal = fault(worn_timing_belt) =====

[Backward-Chain] Q: is 'engine_ticking' present? (auto) -> yes

[Backward-Chain] Q: is 'check_engine_light' present? (auto) -> yes

[Backward-Chain] Q: is 'stalls_after_start' present? (auto) -> yes

GOAL PROVED: fault(worn_timing_belt) confirmed.

Recommendation: Inspect and replace the timing belt immediately - risk of engine damage.

===== BACKWARD CHAINING RUN: goal = fault(electrical_system_failure) =====

[Backward-Chain] Q: is 'dim_headlights' present? (auto) -> yes

[Backward-Chain] Q: is 'slow_engine_crank' present? (auto) -> yes

[Backward-Chain] Q: is 'no_crank' present? (auto) -> yes

[Backward-Chain] Q: is 'battery_light' present? (auto) -> yes

[Backward-Chain] Q: is 'stalls_after_start' present? (auto) -> yes

GOAL PROVED: fault(electrical_system_failure) confirmed.

Recommendation: Have the entire charging/electrical system tested by a technician.

===== ALL TESTS COMPLETE =====

EVALUATION RUBRICS:

Criterion	Weightage (%)	Excellent	Good	Average	Poor
1. Problem Analysis and Domain Understanding	10%	9–10% – Clearly analyzes the real-world problem, domain entities, facts, conditions, goals, and expected conclusions.	7–8% – Good understanding of the domain with minor omissions.	5–6% – Basic domain understanding with several missing elements.	0–4% – Poor or incorrect understanding of the problem domain.
2. Knowledge Acquisition and Representation	10%	9–10% – Acquires relevant domain knowledge and represents it accurately using well-structured Prolog facts and rules.	7–8% – Knowledge is appropriately represented with minor gaps.	5–6% – Basic knowledge representation with missing or unclear facts/rules.	0–4% – Inadequate or inappropriate knowledge representation.
3. Production Rule / Knowledge Base Design	15%	13–15% – Develops a comprehensive, consistent, and logically structured production-rule knowledge base with appropriate facts and rules.	10–12% – Well-designed knowledge base with minor inconsistencies or omissions.	7–9% – Basic rule base developed but several rules are incomplete or unclear.	0–6% – Knowledge base is incomplete, inconsistent, or inappropriate.
4. Prolog Implementation	15%	13–15% – Implements a fully functional expert system using Prolog facts, predicates, rules, variables, unification, and backtracking effectively.	10–12% – Functional Prolog implementation with minor technical issues.	7–9% – Basic implementation works with limited use of Prolog features.	0–6% – Implementation is incomplete, incorrect, or non-functional.

5. Forward Chaining	10%	9–10% – Correctly implements and demonstrates forward chaining from initial facts through applicable rules to final conclusions.	7–8% – Forward chaining works correctly with minor limitations.	5–6% – Basic forward reasoning demonstrated but with incomplete rule processing.	0–4% – Forward chaining is missing or incorrectly implemented.
6. Backward Chaining	10%	9–10% – Effectively demonstrates goal-driven backward chaining using Prolog's inference mechanism and clearly traces the reasoning process.	7–8% – Backward chaining works correctly with minor gaps.	5–6% – Basic backward reasoning demonstrated with limited explanation.	0–4% – Backward chaining is missing or incorrectly implemented.
7. Procedural and Non-Procedural Paradigm Analysis	10%	9–10% – Clearly distinguishes procedural and declarative/non-procedural approaches and demonstrates their application in the Prolog system.	7–8% – Provides a good distinction with minor conceptual gaps.	5–6% – Basic distinction is explained but lacks practical demonstration.	0–4% – Paradigms are misunderstood or not addressed.
8. Inference, Unification and Reasoning Accuracy	10%	9–10% – Produces accurate conclusions using unification, backtracking, and logical inference across diverse test cases.	7–8% – Reasoning is mostly accurate with minor errors.	5–6% – Basic inference works but produces occasional incorrect results.	0–4% – Inference is unreliable or incorrect.

9. Testing, Results and System Demonstration	5%	5% – Provides comprehensive test cases, accurate outputs, reasoning traces, and a convincing live demonstration.	4% – Good testing and demonstration with minor gaps.	2–3% – Limited test cases and basic demonstration.	0–1% – Testing or demonstration is missing/inadequate.
10. Documentation, GitHub and Technical Presentation	5%	5% – Complete documentation, well-organized GitHub repository, clear code, references, and professional presentation.	4% – Good documentation and repository with minor issues.	2–3% – Basic documentation/repository with several omissions.	0–1% – Poor documentation, missing repository, or inadequate presentation.
Total	100%				